

# digit-recognizer-1

November 18, 2023

```
[ ]: import tensorflow as tf
      from tensorflow.keras import datasets
      import matplotlib.pyplot as plt
      import torch
      from torchvision import datasets, transforms
      from torch.utils.data import DataLoader, TensorDataset
      from sklearn.model_selection import train_test_split
      from sklearn.utils import shuffle
      import torch.optim as optim
      import torch.nn as nn
      from torch.optim import SGD
      from sklearn.metrics import accuracy_score
```

This is the function that will enable us to print the image and its label of index(index)

```
[ ]: def show_image(images, labels, index):
      image = images[index].view(28, 28) # Reshape the flattened image
      label = labels[index].item()
      figsize=(2, 2)
      plt.figure(figsize=figsize)
      plt.imshow(image, cmap='gray')
      plt.title(f"Label: {label}")
      plt.axis('off')
      plt.show()
```

Here we will define the transform that will convert images to .tensors instead of ndarray, also it normalizes the values to [-1,1]

```
[ ]: # Define a transform to normalize the data
      transform = transforms.Compose([
          transforms.ToTensor(), # Converts PIL Image or numpy.ndarray to torch.
          ↪ Tensor
          transforms.Normalize((0.5,), (0.5,)) # Normalize to range [-1, 1]
      ])
```

In this cell we will download the data and use the transform we init in the last cell also we will split the images and the labels then we will split the training data into validation (0.2) and training (0.8)

```
[ ]: # Download and load the training data
train_dataset = datasets.MNIST(root='./data', train=True, transform=transform,
    ↪download=True)

# Download and load the test data
test_dataset = datasets.MNIST(root='./data', train=False, transform=transform,
    ↪download=True)

images, labels = train_dataset.data.float() , train_dataset.targets.long()


# images, labels = train_dataset.data, train_dataset.targets
test_images, test_labels = test_dataset.data.float(), test_dataset.targets.
    ↪long()

train_images, val_images, train_labels, val_labels = train_test_split(
    images, labels, test_size=0.2, random_state=42, stratify=labels)

train_images, train_labels = shuffle(train_images, train_labels,
    ↪random_state=42)
val_images, val_labels = shuffle(val_images, val_labels, random_state=42)

[ ]: show_image(val_images, val_labels, 10)
```

Label: 6



Here is the main class that does all the work: input\_size, output\_size, hidden\_layers, activation function dropout\_prob

```
[ ]: class FlexibleNN(nn.Module):
```

```

def __init__(self, input_size, output_size, hidden_layers=[128, 64],
activation=nn.ReLU(), dropout_prob=0.5):
    super(FlexibleNN, self).__init__()
    layers = []
    for i in range(len(hidden_layers) - 1):
        layers.extend([
            nn.Linear(hidden_layers[i], hidden_layers[i + 1]),
            nn.LayerNorm(hidden_layers[i + 1]),
            activation,
            nn.Dropout(p=dropout_prob)
        ])
    # Input layer
    layers.insert(0, nn.Linear(input_size, hidden_layers[0]))

    # Output layer
    layers.append(nn.Linear(hidden_layers[-1], output_size))

    self.model = nn.Sequential(*layers)

def forward(self, x):
    return self.model(x)

```

here is the custom training loop that trains the model using sgd

```

[ ]: def train_model(model, train_loader, criterion, optimizer):

    total_loss = 0.0
    total_samples=0
    correct_predictions=0

    for batch in train_loader:
        images, labels = batch

        # Flatten the images
        images = images.view(images.size(0), -1)

        # Zero the gradients
        optimizer.zero_grad()

        # Forward pass
        outputs = model(images)

        # Calculate the loss
        loss = criterion(outputs, labels)

        # Backward pass
        loss.backward()

```

```

        # Update weights
        optimizer.step()

        # Accumulate the total loss for the epoch
        total_loss+=loss.item()
        # Calculate accuracy
        _, predicted = torch.max(outputs.data, 1)
        total_samples += labels.size(0)
        correct_predictions += (predicted == labels).sum().item()

    accuracy = correct_predictions / total_samples
    return total_loss / len(train_loader),accuracy

```

This is the validate fn that tests the total loss and accuracy of a provided data

```

[ ]: def validate(model, val_loader, criterion):
    model.eval()
    total_loss = 0.0
    all_preds = []
    all_labels = []
    with torch.no_grad():
        for inputs, labels in val_loader:
            inputs = inputs.view(inputs.size(0), -1)
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            total_loss += loss.item()
            _, preds = torch.max(outputs, 1)
            all_preds.extend(preds.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())
    accuracy = accuracy_score(all_labels, all_preds)

    return total_loss / len(val_loader), accuracy

```

here we choose all our hyperparameters and created tendordatasets from our images/labels also we defined our criterion that specifies the loss

```

[ ]: input_size = 28 * 28
    output_size = 10
    hidden_layers = [128, 64, 32]
    learning_rate = 0.01
    num_of_epochs=10

    # Create a TensorDataset
    dataset = TensorDataset(train_images, train_labels)
    val_data = TensorDataset(val_images, val_labels)
    test_data = TensorDataset(test_images, test_labels)
    # Specify batch size

```

```

batch_size = 64

# Create a DataLoader
train_loader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_data, batch_size=64, shuffle=False)
test_loader = DataLoader(test_data, batch_size=64, shuffle=False)
model = FlexibleNN(input_size, output_size, hidden_layers)
print(model)
# Define the criterion
criterion = nn.CrossEntropyLoss()

# Create an SGD optimizer
optimizer = SGD(model.parameters(), lr=learning_rate)

```

```

FlexibleNN(
  (model): Sequential(
    (0): Linear(in_features=784, out_features=128, bias=True)
    (1): Linear(in_features=128, out_features=64, bias=True)
    (2): LayerNorm((64,)), eps=1e-05, elementwise_affine=True)
    (3): ReLU()
    (4): Dropout(p=0.5, inplace=False)
    (5): Linear(in_features=64, out_features=32, bias=True)
    (6): LayerNorm((32,)), eps=1e-05, elementwise_affine=True)
    (7): ReLU()
    (8): Dropout(p=0.5, inplace=False)
    (9): Linear(in_features=32, out_features=10, bias=True)
  )
)

```

here is the real training that happens

```

[ ]: train_losses = []
    val_losses = []
    train_accuracies = []
    val_accuracies = []
    # Train the model
    for epoch in range(num_of_epochs):
        train_loss, train_accuracy = train_model(model, train_loader, criterion,
        ↪optimizer)
        val_loss, val_accuracy = validate(model, val_loader, criterion)
        train_accuracies.append(train_accuracy)
        train_losses.append(train_loss)
        val_losses.append(val_loss)
        val_accuracies.append(val_accuracy)

        print("Epoch {}/ {}: Train Loss: {:.4f}, Validation Loss: {:.4f},
        ↪Validation Accuracy: {:.4f}".format(
            epoch + 1, num_of_epochs, train_loss, val_loss, val_accuracy))

```

```
# train_model(model, train_loader, criterion, optimizer,
↳ num_epochs=num_of_epochs)
```

```
Epoch 1/10: Train Loss: 1.4143, Validation Loss: 0.5956, Validation Accuracy: 0.8719
Epoch 2/10: Train Loss: 0.3835, Validation Loss: 0.2732, Validation Accuracy: 0.9285
Epoch 3/10: Train Loss: 0.2340, Validation Loss: 0.2070, Validation Accuracy: 0.9427
Epoch 4/10: Train Loss: 0.1732, Validation Loss: 0.1654, Validation Accuracy: 0.9544
Epoch 5/10: Train Loss: 0.1393, Validation Loss: 0.1455, Validation Accuracy: 0.9589
Epoch 6/10: Train Loss: 0.1167, Validation Loss: 0.1261, Validation Accuracy: 0.9633
Epoch 7/10: Train Loss: 0.0995, Validation Loss: 0.1197, Validation Accuracy: 0.9660
Epoch 8/10: Train Loss: 0.0875, Validation Loss: 0.1119, Validation Accuracy: 0.9682
Epoch 9/10: Train Loss: 0.0769, Validation Loss: 0.1073, Validation Accuracy: 0.9690
Epoch 10/10: Train Loss: 0.0699, Validation Loss: 0.1039, Validation Accuracy: 0.9695
```

now we will test our model on test set to get final scores

```
[ ]: test_loss, test_accuracy = validate(model, test_loader, criterion)
      print(test_loss)
      print(test_accuracy)
```

```
0.09942378488274374
0.9693
```

A fn to plot:

```
[ ]: def Plotting (train_losses,train_accuracies,val_losses,val_accuracies,variant):
      # Plotting training and validation loss
      plt.figure(figsize=(10, 5))
      plt.subplot(1, 2, 1)
      plt.plot(train_losses, label='Training Loss')
      plt.plot(val_losses, label='Validation Loss')
      plt.xlabel(variant)
      plt.ylabel('Loss')
      plt.title('Training and Validation Loss')
      plt.legend()


      # Plotting training and validation accuracy
      plt.subplot(1, 2, 2)
      plt.plot(train_accuracies, label='Training Accuracy')
```

```

plt.plot(val_accuracies, label='Validation Accuracy')
plt.xlabel('variant')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.tight_layout()
plt.show()

```

Plotting(train\_losses=train\_losses,train\_accuracies=train\_accuracies,val\_losses=val\_losses,val



as we se as epochs pass our model is learning more and the loss is decreasing for both train and val sets also accuracy is increasing for both

here we will validate the learning rate by trying different rates and catching the acc and loss to plot later

```

[ ]: learning_rates = [0.1, 0.01, 0.001, 0.0001, 0.00001]
all_train_losses = []
all_train_accuracies = []
all_val_losses = []
all_val_accuracies = []

for lr in learning_rates:
    model = FlexibleNN(input_size, output_size, hidden_layers)
    optimizer = SGD(model.parameters(), lr=lr)
    train_losses = []
    val_losses = []
    train_accuracies = []

```

```

val_accuracies = []
for epoch in range(num_of_epochs):
    train_loss, train_accuracy = train_model(model, train_loader, criterion,
    optimizer)
    val_loss, val_accuracy = validate(model, val_loader, criterion)
    train_losses.append(train_loss)
    train_accuracies.append(train_accuracy)
    val_losses.append(val_loss)
    val_accuracies.append(val_accuracy)

    print("Epoch {}/{} - Learning Rate: {:.5f}: Train Loss: {:.4f},
    Validation Loss: {:.4f}, Validation Accuracy: {:.4f}".format(
        epoch + 1, num_of_epochs, lr, train_loss, val_loss, val_accuracy))
    # Append results for the current learning rate to the overall
    lists
    all_train_losses.append(train_losses)
    all_train_accuracies.append(train_accuracies)
    all_val_losses.append(val_losses)
    all_val_accuracies.append(val_accuracies)

```

```

Epoch 1/10 - Learning Rate: 0.10000: Train Loss: 0.9273, Validation Loss:
0.3711, Validation Accuracy: 0.8923
Epoch 2/10 - Learning Rate: 0.10000: Train Loss: 0.2427, Validation Loss:
0.1953, Validation Accuracy: 0.9387
Epoch 3/10 - Learning Rate: 0.10000: Train Loss: 0.1652, Validation Loss:
0.1466, Validation Accuracy: 0.9553
Epoch 4/10 - Learning Rate: 0.10000: Train Loss: 0.1316, Validation Loss:
0.1255, Validation Accuracy: 0.9633
Epoch 5/10 - Learning Rate: 0.10000: Train Loss: 0.1096, Validation Loss:
0.1304, Validation Accuracy: 0.9612
Epoch 6/10 - Learning Rate: 0.10000: Train Loss: 0.0939, Validation Loss:
0.1125, Validation Accuracy: 0.9665
Epoch 7/10 - Learning Rate: 0.10000: Train Loss: 0.0822, Validation Loss:
0.1113, Validation Accuracy: 0.9684
Epoch 8/10 - Learning Rate: 0.10000: Train Loss: 0.0744, Validation Loss:
0.0969, Validation Accuracy: 0.9717
Epoch 9/10 - Learning Rate: 0.10000: Train Loss: 0.0661, Validation Loss:
0.1011, Validation Accuracy: 0.9711
Epoch 10/10 - Learning Rate: 0.10000: Train Loss: 0.0603, Validation Loss:
0.0989, Validation Accuracy: 0.9711
Epoch 1/10 - Learning Rate: 0.01000: Train Loss: 1.4183, Validation Loss:
0.5664, Validation Accuracy: 0.8849
Epoch 2/10 - Learning Rate: 0.01000: Train Loss: 0.3719, Validation Loss:
0.2642, Validation Accuracy: 0.9327
Epoch 3/10 - Learning Rate: 0.01000: Train Loss: 0.2304, Validation Loss:

```



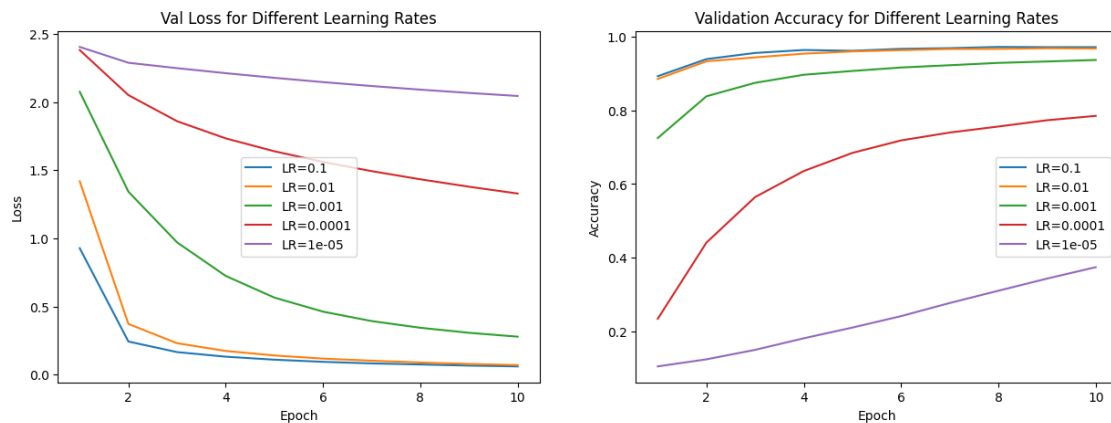
0.2036, Validation Accuracy: 0.9433  
 Epoch 4/10 - Learning Rate: 0.01000: Train Loss: 0.1736, Validation Loss:  
 0.1653, Validation Accuracy: 0.9534  
 Epoch 5/10 - Learning Rate: 0.01000: Train Loss: 0.1406, Validation Loss:  
 0.1477, Validation Accuracy: 0.9595  
 Epoch 6/10 - Learning Rate: 0.01000: Train Loss: 0.1174, Validation Loss:  
 0.1308, Validation Accuracy: 0.9628  
 Epoch 7/10 - Learning Rate: 0.01000: Train Loss: 0.1021, Validation Loss:  
 0.1233, Validation Accuracy: 0.9660  
 Epoch 8/10 - Learning Rate: 0.01000: Train Loss: 0.0890, Validation Loss:  
 0.1180, Validation Accuracy: 0.9663  
 Epoch 9/10 - Learning Rate: 0.01000: Train Loss: 0.0782, Validation Loss:  
 0.1144, Validation Accuracy: 0.9680  
 Epoch 10/10 - Learning Rate: 0.01000: Train Loss: 0.0700, Validation Loss:  
 0.1117, Validation Accuracy: 0.9673  
 Epoch 1/10 - Learning Rate: 0.00100: Train Loss: 2.0755, Validation Loss:  
 1.6099, Validation Accuracy: 0.7248  
 Epoch 2/10 - Learning Rate: 0.00100: Train Loss: 1.3425, Validation Loss:  
 1.1307, Validation Accuracy: 0.8379  
 Epoch 3/10 - Learning Rate: 0.00100: Train Loss: 0.9706, Validation Loss:  
 0.8315, Validation Accuracy: 0.8743  
 Epoch 4/10 - Learning Rate: 0.00100: Train Loss: 0.7243, Validation Loss:  
 0.6348, Validation Accuracy: 0.8962  
 Epoch 5/10 - Learning Rate: 0.00100: Train Loss: 0.5652, Validation Loss:  
 0.5085, Validation Accuracy: 0.9065  
 Epoch 6/10 - Learning Rate: 0.00100: Train Loss: 0.4622, Validation Loss:  
 0.4273, Validation Accuracy: 0.9159  
 Epoch 7/10 - Learning Rate: 0.00100: Train Loss: 0.3929, Validation Loss:  
 0.3692, Validation Accuracy: 0.9221  
 Epoch 8/10 - Learning Rate: 0.00100: Train Loss: 0.3437, Validation Loss:  
 0.3286, Validation Accuracy: 0.9284  
 Epoch 9/10 - Learning Rate: 0.00100: Train Loss: 0.3070, Validation Loss:  
 0.2980, Validation Accuracy: 0.9323  
 Epoch 10/10 - Learning Rate: 0.00100: Train Loss: 0.2784, Validation Loss:  
 0.2742, Validation Accuracy: 0.9363  
 Epoch 1/10 - Learning Rate: 0.00010: Train Loss: 2.3816, Validation Loss:  
 2.1749, Validation Accuracy: 0.2347  
 Epoch 2/10 - Learning Rate: 0.00010: Train Loss: 2.0512, Validation Loss:  
 1.9393, Validation Accuracy: 0.4413  
 Epoch 3/10 - Learning Rate: 0.00010: Train Loss: 1.8595, Validation Loss:  
 1.7886, Validation Accuracy: 0.5647  
 Epoch 4/10 - Learning Rate: 0.00010: Train Loss: 1.7337, Validation Loss:  
 1.6820, Validation Accuracy: 0.6352  
 Epoch 5/10 - Learning Rate: 0.00010: Train Loss: 1.6390, Validation Loss:  
 1.5969, Validation Accuracy: 0.6844  
 Epoch 6/10 - Learning Rate: 0.00010: Train Loss: 1.5608, Validation Loss:  
 1.5246, Validation Accuracy: 0.7183  
 Epoch 7/10 - Learning Rate: 0.00010: Train Loss: 1.4933, Validation Loss:

1.4611, Validation Accuracy: 0.7396  
Epoch 8/10 - Learning Rate: 0.00010: Train Loss: 1.4333, Validation Loss: 1.4041, Validation Accuracy: 0.7559  
Epoch 9/10 - Learning Rate: 0.00010: Train Loss: 1.3789, Validation Loss: 1.3520, Validation Accuracy: 0.7729  
Epoch 10/10 - Learning Rate: 0.00010: Train Loss: 1.3289, Validation Loss: 1.3039, Validation Accuracy: 0.7847  
Epoch 1/10 - Learning Rate: 0.00001: Train Loss: 2.4047, Validation Loss: 2.3106, Validation Accuracy: 0.1057  
Epoch 2/10 - Learning Rate: 0.00001: Train Loss: 2.2884, Validation Loss: 2.2700, Validation Accuracy: 0.1248  
Epoch 3/10 - Learning Rate: 0.00001: Train Loss: 2.2490, Validation Loss: 2.2318, Validation Accuracy: 0.1505  
Epoch 4/10 - Learning Rate: 0.00001: Train Loss: 2.2121, Validation Loss: 2.1964, Validation Accuracy: 0.1818  
Epoch 5/10 - Learning Rate: 0.00001: Train Loss: 2.1780, Validation Loss: 2.1635, Validation Accuracy: 0.2109  
Epoch 6/10 - Learning Rate: 0.00001: Train Loss: 2.1466, Validation Loss: 2.1333, Validation Accuracy: 0.2419  
Epoch 7/10 - Learning Rate: 0.00001: Train Loss: 2.1177, Validation Loss: 2.1056, Validation Accuracy: 0.2775  
Epoch 8/10 - Learning Rate: 0.00001: Train Loss: 2.0913, Validation Loss: 2.0803, Validation Accuracy: 0.3108  
Epoch 9/10 - Learning Rate: 0.00001: Train Loss: 2.0670, Validation Loss: 2.0571, Validation Accuracy: 0.3437  
Epoch 10/10 - Learning Rate: 0.00001: Train Loss: 2.0447, Validation Loss: 2.0357, Validation Accuracy: 0.3746

```
[ ]: plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)
for i, lr in enumerate(learning_rates):
    plt.plot(range(1, num_of_epochs + 1), all_train_losses[i], label=f'LR={lr}')
plt.title('Val Loss for Different Learning Rates')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

# Plotting training and validation accuracy
plt.subplot(1, 2, 2)
for i, lr in enumerate(learning_rates):
    plt.plot(range(1, num_of_epochs + 1), all_val_accuracies[i],
             label=f'LR={lr}')
plt.title('Validation Accuracy for Different Learning Rates')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
```

```
plt.show()
```



### -Higher Learning Rates (e.g., 0.1):

Training Loss: May decrease rapidly initially but may oscillate or diverge. Validation Accuracy: May fluctuate, and the model might not converge to an optimal solution. Overfitting: Higher risk of divergence and poor generalization. Computational Efficiency: Faster initial updates but may require careful tuning. Convergence Speed: Faster initial convergence, but stability is a concern.

### -Moderate Learning Rates (e.g., 0.0001):

Training Loss: Decreases steadily without oscillations or divergence. Validation Accuracy: Shows stable improvement, and the model converges well. Overfitting: Moderate risk; better generalization compared to high learning rates. Computational Efficiency: Efficient updates with a balance between speed and stability. Convergence Speed: Moderate convergence speed with good stability.

### -Lower Learning Rates (e.g., 0.000001):

Training Loss: Decreases more slowly but steadily. Validation Accuracy: Improves steadily with a higher likelihood of convergence. Overfitting: Lower risk, tends to generalize well. Computational Efficiency: Slower updates, but stability is prioritized. Convergence Speed: Slower but more stable convergence.

as we can see the optimal lr is 0.1

now we will try different batch sizes

```
[ ]: batch_sizes = [32, 64, 128, 256, 512]
all_train_losses = []
all_train_accuracies = []
all_val_losses = []
all_val_accuracies = []

for batch_size in batch_sizes:
    model = FlexibleNN(input_size, output_size, hidden_layers)
    train_loader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
```

```

val_loader = DataLoader(val_data, batch_size=batch_size, shuffle=False)
optimizer = SGD(model.parameters(), lr=learning_rate)
train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []

for epoch in range(num_of_epochs):
    train_loss, train_accuracy = train_model(model, train_loader,
↪criterion, optimizer)
    val_loss, val_accuracy = validate(model, val_loader, criterion)
    train_losses.append(train_loss)
    train_accuracies.append(train_accuracy)
    val_losses.append(val_loss)
    val_accuracies.append(val_accuracy)

    print("Epoch {}/{} - Batch Size: {}: Train Loss: {:.4f}, Validation_
↪Loss: {:.4f}, Validation Accuracy: {:.4f}".format(
        epoch + 1, num_of_epochs, batch_size, train_loss, val_loss,
↪val_accuracy))
    all_train_losses.append(train_losses)
    all_train_accuracies.append(train_accuracies)
    all_val_losses.append(val_losses)
    all_val_accuracies.append(val_accuracies)

```

```

Epoch 1/10 - Batch Size: 32: Train Loss: 1.1597, Validation Loss: 0.4261,
Validation Accuracy: 0.8963
Epoch 2/10 - Batch Size: 32: Train Loss: 0.2819, Validation Loss: 0.2081,
Validation Accuracy: 0.9403
Epoch 3/10 - Batch Size: 32: Train Loss: 0.1765, Validation Loss: 0.1543,
Validation Accuracy: 0.9559
Epoch 4/10 - Batch Size: 32: Train Loss: 0.1351, Validation Loss: 0.1403,
Validation Accuracy: 0.9611
Epoch 5/10 - Batch Size: 32: Train Loss: 0.1089, Validation Loss: 0.1251,
Validation Accuracy: 0.9638
Epoch 6/10 - Batch Size: 32: Train Loss: 0.0935, Validation Loss: 0.1142,
Validation Accuracy: 0.9666
Epoch 7/10 - Batch Size: 32: Train Loss: 0.0809, Validation Loss: 0.1064,
Validation Accuracy: 0.9689
Epoch 8/10 - Batch Size: 32: Train Loss: 0.0722, Validation Loss: 0.1056,
Validation Accuracy: 0.9680
Epoch 9/10 - Batch Size: 32: Train Loss: 0.0640, Validation Loss: 0.0967,
Validation Accuracy: 0.9712
Epoch 10/10 - Batch Size: 32: Train Loss: 0.0574, Validation Loss: 0.1016,
Validation Accuracy: 0.9699
Epoch 1/10 - Batch Size: 64: Train Loss: 1.4080, Validation Loss: 0.5613,
Validation Accuracy: 0.8877

```

Epoch 2/10 - Batch Size: 64: Train Loss: 0.3624, Validation Loss: 0.2668,  
Validation Accuracy: 0.9301  
Epoch 3/10 - Batch Size: 64: Train Loss: 0.2244, Validation Loss: 0.1958,  
Validation Accuracy: 0.9462  
Epoch 4/10 - Batch Size: 64: Train Loss: 0.1687, Validation Loss: 0.1629,  
Validation Accuracy: 0.9543  
Epoch 5/10 - Batch Size: 64: Train Loss: 0.1335, Validation Loss: 0.1404,  
Validation Accuracy: 0.9609  
Epoch 6/10 - Batch Size: 64: Train Loss: 0.1107, Validation Loss: 0.1320,  
Validation Accuracy: 0.9610  
Epoch 7/10 - Batch Size: 64: Train Loss: 0.0967, Validation Loss: 0.1212,  
Validation Accuracy: 0.9647  
Epoch 8/10 - Batch Size: 64: Train Loss: 0.0842, Validation Loss: 0.1094,  
Validation Accuracy: 0.9687  
Epoch 9/10 - Batch Size: 64: Train Loss: 0.0752, Validation Loss: 0.1139,  
Validation Accuracy: 0.9670  
Epoch 10/10 - Batch Size: 64: Train Loss: 0.0672, Validation Loss: 0.1114,  
Validation Accuracy: 0.9673  
Epoch 1/10 - Batch Size: 128: Train Loss: 1.6558, Validation Loss: 0.8308,  
Validation Accuracy: 0.8538  
Epoch 2/10 - Batch Size: 128: Train Loss: 0.5376, Validation Loss: 0.3770,  
Validation Accuracy: 0.9151  
Epoch 3/10 - Batch Size: 128: Train Loss: 0.3079, Validation Loss: 0.2680,  
Validation Accuracy: 0.9337  
Epoch 4/10 - Batch Size: 128: Train Loss: 0.2291, Validation Loss: 0.2060,  
Validation Accuracy: 0.9469  
Epoch 5/10 - Batch Size: 128: Train Loss: 0.1839, Validation Loss: 0.1738,  
Validation Accuracy: 0.9532  
Epoch 6/10 - Batch Size: 128: Train Loss: 0.1541, Validation Loss: 0.1633,  
Validation Accuracy: 0.9560  
Epoch 7/10 - Batch Size: 128: Train Loss: 0.1342, Validation Loss: 0.1507,  
Validation Accuracy: 0.9570  
Epoch 8/10 - Batch Size: 128: Train Loss: 0.1175, Validation Loss: 0.1359,  
Validation Accuracy: 0.9623  
Epoch 9/10 - Batch Size: 128: Train Loss: 0.1062, Validation Loss: 0.1282,  
Validation Accuracy: 0.9646  
Epoch 10/10 - Batch Size: 128: Train Loss: 0.0961, Validation Loss: 0.1223,  
Validation Accuracy: 0.9659  
Epoch 1/10 - Batch Size: 256: Train Loss: 1.8998, Validation Loss: 1.2363,  
Validation Accuracy: 0.8145  
Epoch 2/10 - Batch Size: 256: Train Loss: 0.8813, Validation Loss: 0.6439,  
Validation Accuracy: 0.8893  
Epoch 3/10 - Batch Size: 256: Train Loss: 0.5151, Validation Loss: 0.4292,  
Validation Accuracy: 0.9107  
Epoch 4/10 - Batch Size: 256: Train Loss: 0.3671, Validation Loss: 0.3250,  
Validation Accuracy: 0.9266  
Epoch 5/10 - Batch Size: 256: Train Loss: 0.2921, Validation Loss: 0.2748,  
Validation Accuracy: 0.9326


Epoch 6/10 - Batch Size: 256: Train Loss: 0.2452, Validation Loss: 0.2348, Validation Accuracy: 0.9403  
Epoch 7/10 - Batch Size: 256: Train Loss: 0.2114, Validation Loss: 0.2249, Validation Accuracy: 0.9417  
Epoch 8/10 - Batch Size: 256: Train Loss: 0.1867, Validation Loss: 0.1942, Validation Accuracy: 0.9505  
Epoch 9/10 - Batch Size: 256: Train Loss: 0.1675, Validation Loss: 0.1735, Validation Accuracy: 0.9557  
Epoch 10/10 - Batch Size: 256: Train Loss: 0.1512, Validation Loss: 0.1702, Validation Accuracy: 0.9559  
Epoch 1/10 - Batch Size: 512: Train Loss: 2.0277, Validation Loss: 1.5002, Validation Accuracy: 0.7015  
Epoch 2/10 - Batch Size: 512: Train Loss: 1.2016, Validation Loss: 0.9758, Validation Accuracy: 0.8516  
Epoch 3/10 - Batch Size: 512: Train Loss: 0.8251, Validation Loss: 0.6993, Validation Accuracy: 0.8847  
Epoch 4/10 - Batch Size: 512: Train Loss: 0.6087, Validation Loss: 0.5349, Validation Accuracy: 0.9022  
Epoch 5/10 - Batch Size: 512: Train Loss: 0.4783, Validation Loss: 0.4382, Validation Accuracy: 0.9117  
Epoch 6/10 - Batch Size: 512: Train Loss: 0.3957, Validation Loss: 0.3670, Validation Accuracy: 0.9237  
Epoch 7/10 - Batch Size: 512: Train Loss: 0.3391, Validation Loss: 0.3263, Validation Accuracy: 0.9269  
Epoch 8/10 - Batch Size: 512: Train Loss: 0.2993, Validation Loss: 0.2872, Validation Accuracy: 0.9346  
Epoch 9/10 - Batch Size: 512: Train Loss: 0.2682, Validation Loss: 0.2613, Validation Accuracy: 0.9375  
Epoch 10/10 - Batch Size: 512: Train Loss: 0.2443, Validation Loss: 0.2420, Validation Accuracy: 0.9411

```
[ ]: plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)
for i, batch_size in enumerate(batch_sizes):
    plt.plot(range(1, num_of_epochs + 1), all_val_losses[i], label=f'Batch_
    ↳Size={batch_size}')
plt.title('Val Loss for Different Batch Sizes')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

# Plotting training and validation accuracy
plt.subplot(1, 2, 2)
for i, batch_size in enumerate(batch_sizes):
    plt.plot(range(1, num_of_epochs + 1), all_val_accuracies[i], label=f'Batch_
    ↳Size={batch_size}')
plt.title('Validation Accuracy for Different Batch Sizes')
```

```
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



### -Smaller Batch Sizes (e.g., 32, 64):

Training Loss: Decreases more erratically due to noisy updates. Validation Accuracy: May achieve higher accuracy initially but has a risk of overfitting. Overfitting: Higher risk due to quicker adaptation to the training data. Computational Efficiency: Faster updates but potentially less GPU utilization. Convergence Speed: Faster convergence but may converge to a suboptimal solution quickly.

### -Larger Batch Sizes (e.g., 256, 512):

Training Loss: Decreases more steadily, resulting in smoother convergence. Validation Accuracy: May increase more steadily and generalize better. Overfitting: Lower risk due to a more representative sample in each update. Computational Efficiency: More GPU utilization, potentially faster overall training time. Convergence Speed: Slower convergence in terms of updates per epoch but may converge to a more stable solution.